BrandProof

Every claim, accounted for.

A brand-governance workspace that checks marketing captions and posters against owner-confirmed facts, wording and assets, explains the evidence, and leaves the final decision to a person.

Build status: the central workflow works locally. Free hosting is prepared, but a public project URL, remote persistence acceptance and public repository access are still pending. This document is an updated submission draft until those items are completed.

Team and responsibilities

Member	Primary responsibility
Rushab Ramhari Patil	Backend, AI integration, OCR and database
Yogita	React frontend, workflow and integration

Team: YOGI2KANWAR. The team leader must confirm the final submission and member details. Implementation and tests include substantial AI assistance; members should review and understand the submitted work.

WORKING PROJECT URL - ENTER THE VERIFIED PUBLIC HTTPS URL

SOURCE REPOSITORY - CURRENTLY PRIVATE; PUBLIC ACCESS REQUIRED

https://github.com/YogitaKanwar/BrandProof

Prepared branch: **fix/demo-session-resume**. It combines main commit 833c6ff with session recovery and the hosting package. The feature branch has been pushed; main has not been changed by this continuation.

Round 2 requires a working project link, source link, updated PDF, completed features, implementation details, changes from the plan and three substantive Round 3 improvements. Only the team leader submits. These requirements come from the organiser email supplied by the team.

01 / The judge's core journey

The product makes its comparison basis explicit. It cannot know a correct price or date without an approved reference. Missing, conflicting or uncertain evidence stays incomplete rather than becoming a confident verdict.

Step	What the judge can do
1. Confirm the reference	Create a campaign and add typed facts or approved wording. A demo owner explicitly confirms a source before it can be used. The editable example supplies fictional confirmed rules for exploration.
2. Submit content	Save a caption or choose a PNG/JPEG poster. The combined upload-and-review action waits for the actual saved poster revision before reviewing.
3. Read explained issues	See the offending caption span or extracted poster text, rule explanation, immutable source/version evidence and an unapplied proposed correction.
4. Make a human decision	Apply a correction explicitly, discuss it, request review and record a reviewer decision. Incomplete checks require explicit manual notes; they never become automated passes.
5. Preserve accountability	Reload the same workspace in the same browser. Older runs remain historical; later content/source edits make prior readiness stale.

Actual local browser result

The fictional workshop caption produced four issues: INR 599 versus confirmed INR 499; 2026-11-13 versus confirmed 2026-11-12; "cheap" versus "affordable"; and missing "Terms apply". Reload restored the same campaign and stored findings. Contextual AI remained unavailable and the overall review stayed partial.

Finding	Submitted content	Confirmed reference
Price	INR 599	INR 499
Date	2026-11-13	2026-11-12
Terminology	cheap	affordable
Required wording	Missing	Terms apply

Actual outputs from fictional input, independently exercised through the browser and real HTTP smoke tests. The reference is supplied and confirmed by the owner; the system does not invent it.

02 / Completed features

Implemented and exercised locally in the combined code. Hosting and provider availability are separate acceptance conditions.

Capability	Delivered behavior and practical scope
Approved reference records	Six typed rule kinds: price, date, terminology, required text, tone and claim policy. Owner confirmation, pending/approved states, validity and immutable versions. Background notes are not automatically parsed.
Caption checks	Bounded exact price, ISO-date, whole-term and required-wording checks. Ambiguous scope or conflicting references request clarification. This is not arbitrary fact checking.
Poster OCR	PNG/JPEG validation, canonical private image storage, real Tesseract extraction and scaled text boxes. Unavailable, low-confidence and timed-out extraction stay partial.
Logo reference and comparison	Private owner-confirmed logo versions and conservative bounded local template comparison. A close candidate is still for manual checking; raw similarity is not calibrated confidence.
Explained review results	Severity/category, spans/boxes where supported, approved source snapshots, proposed actions and per-check coverage. No findings does not imply approval.
Corrections and discussion	Explicit caption correction, persisted comments, before/current revisions and append-only finding decisions. Conflicts preserve acknowledged server state.
Human review and receipts	Creator requests, reviewer decisions, incomplete-check notes, receipt/history, self-approval guards and readiness invalidation after edits. Roles are simulated demo roles.
Recovery and persistence	Same-browser recovery with a 24-hour HttpOnly cookie; role tokens remain in memory. Database persistence is exercised locally. Remote redeploy acceptance is pending.
Contextual AI adapter	Groq/OpenRouter adapter with server enablement plus per-review consent, bounded approved-text retrieval, strict citation/span validation and persistent admission limits. Hosting configuration defaults off.

Not claimed: general truth verification; production identity; reliable semantic logo compliance; automated typography/contrast/accessibility checks; automatic PDF guideline extraction; a public deployment; or a successful representative model benchmark.

03 / Technology and implementation

Layer	Implementation
Interface	React 19, TypeScript and Vite. One workspace for references, content, findings and human review. Hash routes support review deep links.
API and schema	Python 3.12+, FastAPI, Pydantic and generated OpenAPI. SQLAlchemy with exercised SQLite development storage and a psycopg 3 PostgreSQL hosting path exercised in container CI.
Images and OCR	Pillow validation/canonicalisation, private database image bytes and local Tesseract English OCR. Bounded subprocess output, timeout and cleanup.
Review engine	Deterministic rule checks plus optional consented provider review over selected approved text. Every run snapshots content/source revisions and preserves incomplete coverage.
Human workflow	Role-scoped demo capabilities, optimistic revisions, persisted decisions and revision-bound approval receipts. Demo role switching is not verified identity.
Deployment package	Multi-stage Docker build: Node 24 compiles React; non-root Python/Tesseract runtime serves both UI and API. One worker, PORT support and database health probe.

Prepared free hosting architecture

Browser over HTTPS
React UI and /api/v1 on one Render Free web service

FastAPI + exact checks + Tesseract + human workflow
Optional external model receives caption/selected text only after consent

Neon Free PostgreSQL
Campaigns, source versions, findings, images/logos and receipts persist outside the temporary app filesystem

The Blueprint asks for the database URL privately, enables secure cookies and leaves external AI off. Secrets and local .env files are excluded from the image. Same-origin requests avoid a separate frontend/API hosting split. Ordinary hosted PostgreSQL URLs select the installed driver automatically.

Render Free can sleep after 15 idle minutes and take about one minute to wake. Neon Free currently provides 1 GB/project and 100 CU-hours/project/month; usage limits still apply. This architecture is prepared, not deployed or remotely verified.

Official hosting references, checked 28 October 2026: render.com/docs/free; neon.com/pricing. Setup: docs/FREE_HOSTING.md and render.yaml in the source branch.

04 / Verification and plan changes

Evidence	Actual result
Backend regression	445 passed; no failures or skips. Full combined local run: 39.05 seconds. Ruff lint and formatting passed.
Frontend checks	47 unit tests passed; production build and lint passed. All seven actual HTTP integration scripts passed against an isolated local database.
End-to-end HTTP	Caption smoke: four real issues, zero exact issues after correction, historical run stale. Parser smoke: 10 readback cases. Workflow smoke: actual OCR, one poster contradiction, explicit correction, persisted discussion/decisions, manual-supplement approval, then stale after poster edit.
Packaged browser app	Built UI and API served together locally. Editable example produced four findings. Reload restored the same campaign and review. Full dedicated Chrome/CDP regression was not rerun here.
Container + PostgreSQL CI	GitHub Actions run 40671322353 passed on commit 2e369f1: complete Docker build, UI, actual caption/parser/OCR/human workflows against PostgreSQL 16, plus TLS-origin/secure-cookie recovery and forged-origin rejection. This is a CI test, not a public deployment.
Provider evidence	Prior real Groq calls and persisted output are documented separately. A six-case synthetic evaluation had three assessed development matches and three unassessed heldout cases. That is not representative accuracy. No new provider calls were made for this integration.
Still pending	Actual Render/Neon deployment, public HTTPS and sleep/redeploy persistence, public links and independent human usability/accessibility acceptance. CI does not establish hosted acceptance.

Changes from the Round 1 plan

Reduced for a working core: manual typed, owner-confirmed references replace automatic guideline extraction; bounded keyword retrieval replaces embeddings/vector search; synchronous bounded review replaces a durable job queue; simulated role capabilities replace production membership. Broad poster claims, typography and accessibility remain manual review scope.

Added during the build: conservative reference-logo comparison; one-action poster upload followed by revision-checked review; same-browser session recovery; an editable fictional example; and a single-origin free hosting package. These additions support the central submission-to-human-decision path.

What distinguishes the approach: separate contradiction, missing evidence and unassessed work; retain the exact approved reference version behind a finding; and invalidate readiness after edits. A score or an empty results panel cannot silently certify a campaign.

Evidence: docs/task-status.md, docs/configured-backend-tests.md, docs/LOGO_COMPARISON_REVIEW.md and docs/examples/live-model-evaluation.json. CI: Actions run 40671322353 (repository access required). Automated reviewer actions use fictional data and are not human-user feedback.

05 / Three major Round 3 additions

1. Approved-guideline import with traceable extraction

Import PDF/document guidelines, extract candidate rules with page/section evidence, and require an owner to confirm each candidate before activation. Add scope, conditions and validity to facts such as price and dates. Preserve both the original document and the confirmed structured version.

Acceptance: a judge imports a real guideline, corrects or rejects extracted candidates, confirms a rule and follows a finding back to its exact document location. Unconfirmed extraction can never become a source of truth.

2. Measured visual brand review

Extend beyond the current template candidate to explicit logo placement/clear-space checks, approved colours and supported typography/contrast checks. Build a frozen labelled poster set containing real distortions, low resolution and difficult backgrounds. Report coverage and abstention per visual check.

Acceptance: report false positives, missed defects and abstentions on a heldout set. Each supported visual issue has an actual region and reference. Ambiguous cases stay manual; no uncalibrated similarity is presented as confidence.

3. Real team approval and recurring-error insights

Add authenticated organisation membership, actual owner/creator/reviewer invitations and permissions, a durable review queue, and campaign analytics over recurring confirmed issues. Track dismissal/correction decisions so the team can see which source rules create repeated problems.

Acceptance: separate people complete a review without role switching; unauthorised access is denied; interrupted queued work recovers; and an insights view drills from a recurring issue into its real runs and human decisions.

Final submission handoff

Before the team leader submits: create the free hosting accounts; deploy the prepared branch; verify HTTPS, actual OCR, human decision and reload/redeploy persistence; fill the project URL on page 1; make source/PDF links publicly accessible; and open every link in a private window. Submit only through the organiser's official form.

The supplied organiser email states a hard deadline of 28 October at 10:00 AM IST. This update is prepared on 28 October; it is not evidence of an on-time submission or an extension. Confirm any extension directly with the organisers.