BrandProof

Evidence before every publish. Team YOGI2KANWAR - Codex 48, Round 3. Final acceptance: 2 November 2026.

BrandProof helps a campaign team check marketing copy and posters against confirmed brand references, inspect why a finding was raised, and record an accountable human decision. It keeps content, guideline versions, automated coverage and human approval connected in one review workspace.

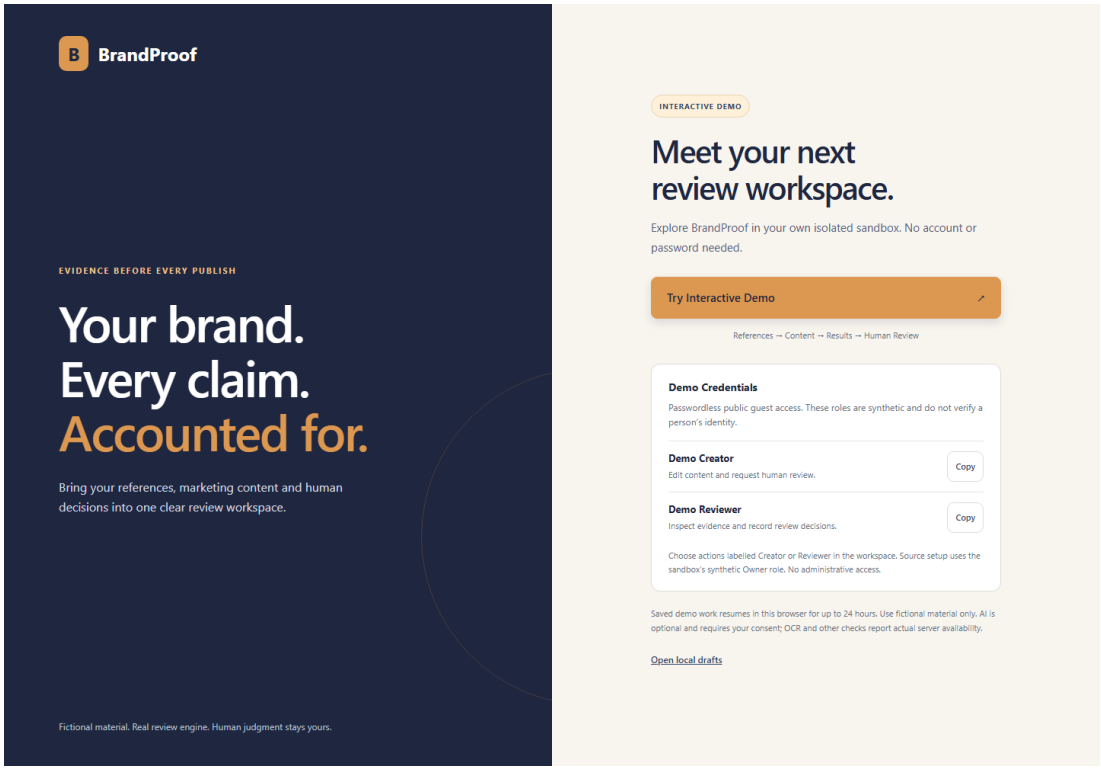
Exactly three promised improvements

1. PDF guideline import with traceable evidence. Import supported directives, inspect exact page quotations, then explicitly confirm versioned rules.

2. Measured visual consistency and accessibility. Inspect supported background colour, contrast and poster-description checks, with measurements and honest uncertainty.

3. Governance Insights. Investigate recurring findings, exact guideline versions and reviewer reasons using persisted evidence and minimum sample requirements.

The existing OCR, human-review and bounded advertising-risk capabilities support these three improvements. They are not additional Round 3 promises.



Genuine isolated release-preview entry at 5eac376. Passwordless access creates an isolated fictional sandbox. Production remains unchanged.

Judge demo: brand-proof-tau.vercel.app/#/demo

Repository: github.com/YogitaKanwar/BrandProof

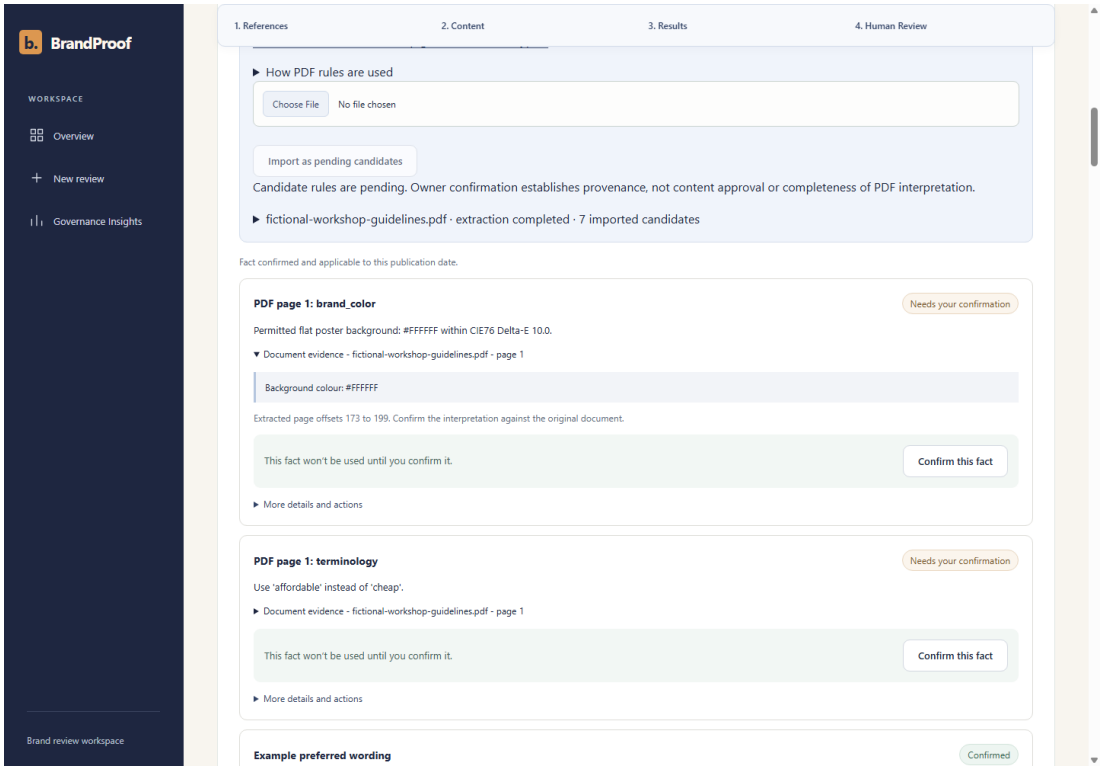
1. Traceable PDF guidelines

Owners import text-layer PDFs up to 5 MiB / 30 pages. A bounded pypdf worker extracts supported directive candidates with page number, exact quotation/offsets and document SHA-256. Candidates are pending; no unconfirmed rule is authoritative evidence.

Explicit synthetic Owner confirmation creates approved v2 from pending v1 while preserving the original page evidence. Immutable source history and saved review snapshots retain provenance. Editing a rule clears its old citation when the wording changes, preventing a corrected rule from misleadingly inheriting an unsupported quote.

What the real workflow demonstrates

The compact Workshop PDF produces three candidates; the extended sample produces seven rule types. Both specify INR 899. A real review of sample marketing content at INR 599 produces a source-version-cited contradiction. Correcting supported caption facts does not approve content or fix an unchanged poster.



Genuine isolated-preview PDF candidate with expanded page quotation and exact provenance. The rule remains pending until explicit Owner confirmation.

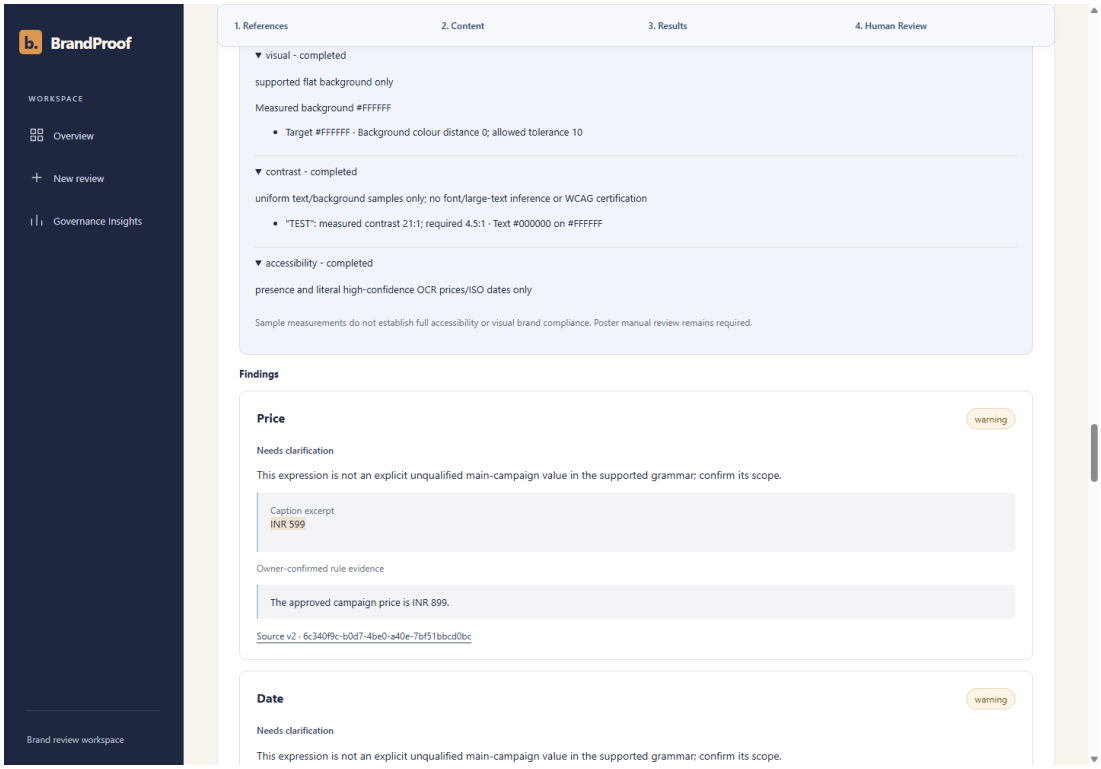
Robustness and limits

Tests exercise both shipped files, two supported pages, mixed supported/blank pages, actual raster-only and blank PDFs, encryption/corruption, wrong type, byte/page limits, extraction timeout, repeated imports, wrong roles, stale revisions and changed-rule provenance. Reimports create separate pending candidates; conflicting approved rules require clarification. Scans and unsupported narrative rules need manual transcription. Import does not authenticate document ownership.

2. Measured checks, stated scope

A reliable flat-background sample exposes the observed hex colour, approved target, CIE76 deviation and source tolerance. Supported OCR-word samples expose measured foreground/background, contrast ratio and the strictest confirmed threshold. Findings cite the relevant confirmed source.

Description checks assess presence and literal high-confidence OCR price/date coverage. Saving a description creates a content revision and can stale a prior approval. Description presence is not a semantic-quality verdict.



Genuine isolated-preview measurements after replacing the Workshop image with a supported controlled black/white TEST poster: contrast 21:1 against 4.5:1. The Workshop image truthfully abstains; limited samples do not establish full accessibility.

Measured acceptance

Real rendered/OCR white-black and white-grey fixtures, approved/mismatched backgrounds, exact contrast thresholds, ordinary antialiasing, uncertainty, missing/whitespace descriptions, permissions and stale receipts passed. The combined public browser workflow passed actual measured checks and description handling.

Gradients, photographs, mixed ink, noisy or small/rotated text and incomplete OCR may be partial or unassessed. At most 100 OCR words are sampled. Eighteen optional Linux-path font-matrix cases were skipped on Windows; required Noto Sans fixture tests ran. No universal visual conformity, typography verification or full WCAG certification is claimed.

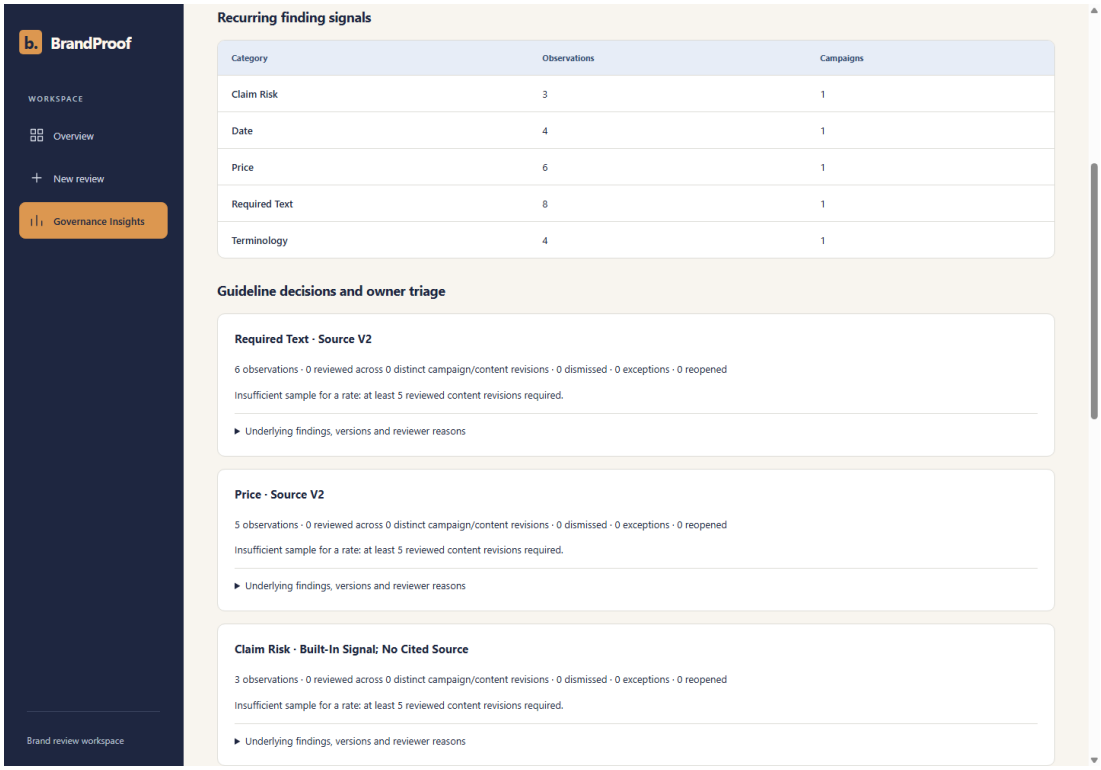
The candidate fixes a reproduced description-loading race after poster replacement. On the older public build, load the saved description before typing; public measurement acceptance uses that recovery action. The candidate passes replacement/edit/save without it.

3. Governance from saved evidence

Insights show recurring categories, campaign counts, exact confirmed source-version groups, latest reviewer dispositions/reasons and links to the underlying saved findings. Deleted campaigns remain hidden and workspaces stay isolated.

Minimum sample and denominator

A group needs at least five reviewed observations and five distinct reviewed campaign/content revisions before a rate is shown. The rate is dismissals plus exceptions divided by all reviewed observations, including reopened decisions. Identical content/source/pack reruns contribute the latest snapshot and do not inherit previous-run decisions.



Genuine isolated-preview Governance Insights. Insufficient sample is visible; small samples do not become numerical trend recommendations.

Verified boundaries

API tests passed zero, one, four, exactly five and six reviewed snapshots, missing-reason rejection, updated/reopened decisions, multiple campaigns/source versions, reruns, run/decision truncation and isolation. Browser acceptance followed an insight into the exact historical run/finding and recovered it after refresh.

The bounded result covers at most 500 runs and 5,000 decisions, with at most 50 evidence links per group and explicit truncation. Signals invite inspection of reasons, not proof a guideline is wrong. Regulatory groups never recommend weakening an obligation because findings were dismissed. All demonstrated decisions are synthetic QA, not real customer analytics.

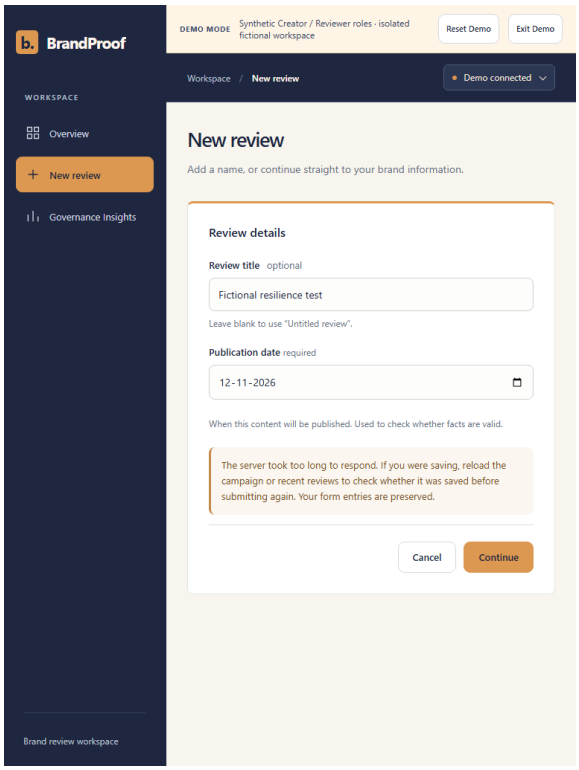
Verification and user handling

Final local backend: 594 collected; **576 passed, 18 optional font skips, zero failures**. **Frontend: 62 passed, zero failures/skips**; lint, TypeScript/Vite build and backend/scripts Ruff checks passed. No code coverage percentage or representative OCR/model accuracy was measured.

Real HTTP: nine frontend suites passed for sessions/campaigns, sources, reviews, posters, collaboration, approval, consent, removal/restore and both PDFs. Four Python workflow scripts passed; parser smoke verified ten readback cases, and real Tesseract workflow verified a reasoned scripted approval becoming stale after poster edit.

Genuine Chromium: three local and three isolated-preview suites passed: judge journey, combined improvements and failure injection. Fresh contexts exercised roles/isolation, OCR boxes, human decisions, PDF confirmation, numeric measurements, Insights links, refresh, reset/logout and removal/Undo. Desktop 1440px, tablet 768px and narrow 390/320px layouts were exercised. Preview description replacement/edit/save passes without manual reload.

Failure injection: a 503 and an actual ten-second transport timeout preserved the draft and restored controls. Duplicate activation sent one real create request. Wrong/empty/oversized/corrupt PDFs failed clearly, then a valid retry created only pending rules. New error-message regressions prevent PDF failures from prescribing image-only remedies.



Genuine isolated-preview tablet screenshot after an actual ten-second transport timeout: input remains available and recovery asks users to check saved state before resubmitting.

Actual local process restart preserved ten compared resources, the recovery workspace, credentials and all role access. Clean isolated Windows setup installed locked dependencies, passed frontend lint/62 tests/build, read-only migration preflight and restart persistence. Linux instructions follow the existing Docker/CI path; they were not executed on this host.

Release, limits and contributions

Current public release and candidate

Production Vercel and Render remain **e01f52b6ee02b6c85c480619cfc82cc486a56123**. Candidate **5eac37660a7c4f53d17464dea3ced563486aa9bd** was merged by Yogita through PR 14, creating main **1763f2e**. Release evidence follows in PR 15. No production deployment or migration was performed; production approval remains separate.

The isolated full-stack preview runs the candidate with its own disposable SQLite database, AI disabled and automatic deployment off. It passed all three browser suites, both-file PDF HTTP checks and measured-feature smoke. Later documentation commits preserve the tested application tree.

Public startup initially returned a proxy 502; direct Render database health recovered in about 22 seconds. The initial browser harness stopped at 45 seconds, shorter than the app's documented 90-second readiness. The harness now allows that initial window; unchanged assertions and the warmed public judge/PDF/combined journeys passed. Both attempts are retained. This is not a guarantee of idle cold-start latency.

Exact-candidate push CI and PR CI passed Docker/PostgreSQL 16, six-step legacy upgrades, real workflows, 576 backend tests (18 optional skips) and 62 frontend tests/lint/build. Backend/schema code matches live; **no new migration is needed**. Recommended release is frontend-only, with the existing API/database retained and the previous Vercel deployment available for frontend rollback. Production remains held for explicit approval.

Remaining limits

Text-layer/directive PDFs only; limited image samples; OCR omissions/logo timeouts; bounded English/India legal scope; contextual AI unavailable. Synthetic actors are not verified identities. No full WCAG/brand/legal certification, representative accuracy or verified production backup/restore is claimed. Provider metadata confirms free PostgreSQL expiry **27 November 2026**. Direct schema SQL remains blocked by its unchanged external allowlist; frontend-only release avoids a database change.

Technology and actual ownership

React 19, TypeScript, Vite; Python, FastAPI, Pydantic, SQLAlchemy, PostgreSQL/SQLite, pypdf, Pillow and Tesseract; Vercel proxy and Render Docker. Optional external AI remains server-configured and consent-gated.

Team-supplied roles: **Yogita Kanwar - frontend and integration. Rushab Ramhari Patil - backend, OCR, AI and database.** Significant implementation/testing used coding-assistant support. No invented hours, percentages or independent user feedback.

Reproduce using README.md. Exact commands, artifacts, skipped tests and release blockers: docs/ROUND3_FINAL_QA.md. Scope-by-case: docs/ROUND3_REQUIREMENTS_MATRIX.md. User path: docs/JUDGE_FINAL_WALKTHROUGH.md. Disclosure: docs/ROUND3_KNOWN_LIMITATIONS.md. This PDF is generated from docs/ROUND3_SUBMISSION_CONTENT.md. Official submission remains the team's action.